

JobKit - Teradata DBA interview questions

212 questions for Teradata DBA. Each answer is five pointers with working code. Single-file download.

1. What is Teradata?

1. Definition you should say first: A shared-nothing MPP database for large analytic workloads. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Teradata is a shared-nothing MPP analytic database. Vantage is the current platform branding (SQL plus analytics). Each AMP owns a slice of every table.
3. Draw PE -> BYNET -> AMPs. Mention PI, stats, TASM, and BAR in the first minute.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.DiskSpaceV WHERE DatabaseName = 'EDW';  
HELP SESSION;  
SELECT HASHAMP() AS amps;
```
5. Calling it 'just like Oracle but faster' without AMP/PI/skew will not survive a DBA round.

2. What is an AMP?

1. Definition you should say first: Access Module Processor: a parallel worker that stores and processes part of the data. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew). Down AMP triggers fallback/cliquest recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.
3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;  
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;  
-- Viewpoint: AMP usage / skew
```
5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

3. What is a PE?

1. Definition you should say first: Parsing Engine: parses SQL, optimizes, and sends steps to AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PE = Parsing Engine: session handling, parse, optimize, dispatch steps to AMPs.
3. PE CPU spikes on bad dynamic SQL or too many short tactical requests. TASM protects mixed workloads.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT AccountName, UserName FROM DBC.SessionInfoV WHERE UserName = USER;  
EXPLAIN SELECT COUNT(*) FROM edw.fact_orders;
```
5. People confuse PE with AMP. PE does not store the table; AMPs do.

4. What is BYNET?

1. Definition you should say first: The interconnect that moves messages and data between PEs and AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BYNET is the interconnect for messages, row redistribution, and duplication between PEs and AMPs.
3. Product joins and redistributing a huge uncompressed table flood BYNET. Check ResUsage and Viewpoint fabric metrics.
4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN  
SELECT *  
FROM edw.fact_orders f  
JOIN edw.dim_customer d ON d.region = f.region;  
-- look for: redistributing / duplicating all rows of
```
5. If you only add CPU and ignore join geography, BYNET will still stall.

5. What is a vproc?

1. Definition you should say first: A virtual processor such as AMP or PE running on a node. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Cliques = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:  
SELECT * FROM DBC.VprocStatus;
```

JobKit - Teradata DBA interview questions

```
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

6. What is a clique?

1. Definition you should say first: A set of nodes that share disks so a failed node can fail over. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.

3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.

4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
```

```
SELECT * FROM DBC.VprocStatus;
```

```
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

7. What is a hot standby node?

1. Definition you should say first: A spare node that can take over vprocs after a failure. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.

3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.

4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
```

```
SELECT * FROM DBC.VprocStatus;
```

```
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

8. What is Fallback?

1. Definition you should say first: A second copy of a table on another AMP for availability. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Fallback stores a second copy of rows on another AMP so a down AMP does not lose the table.

3. Enable on critical dimensions and control tables. It costs disk and write I/O. Not a substitute for BAR.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_calendar (  
    cal_dt DATE NOT NULL  
)
```

```
PRIMARY INDEX (cal_dt)
```

```
FALLBACK;
```

5. Fallback does not protect against a bad UPDATE. You still need backups and a back-out plan.

9. When do you use Fallback?

1. Definition you should say first: Critical tables that must survive AMP or disk loss. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Fallback stores a second copy of rows on another AMP so a down AMP does not lose the table.

3. Enable on critical dimensions and control tables. It costs disk and write I/O. Not a substitute for BAR.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_calendar (  
    cal_dt DATE NOT NULL  
)
```

```
PRIMARY INDEX (cal_dt)
```

```
FALLBACK;
```

5. Fallback does not protect against a bad UPDATE. You still need backups and a back-out plan.

10. What is a Primary Index?

1. Definition you should say first: The distribution key that decides which AMP stores a row. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.

3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.

4. Working code / syntax (write this on Teradata SQL):

JobKit - Teradata DBA interview questions

```
CREATE MULTISET TABLE stg.orders_in (  
    order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)  
) NO PRIMARY INDEX;  
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE,  
    amt DECIMAL(18,2)  
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

11. Unique vs non-unique PI?

- 1. Definition you should say first: UPI enforces uniqueness; NUPI allows duplicates and can skew. Then spend the rest of the answer on architecture, a working example, and a production failure.**
- 2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.**
- 3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.**

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE stg.orders_in (  
    order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)  
) NO PRIMARY INDEX;  
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE,  
    amt DECIMAL(18,2)  
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

12. What is a Secondary Index?

- 1. Definition you should say first: An extra access path; it has a subtable and extra I/O. Then spend the rest of the answer on architecture, a working example, and a production failure.**
- 2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.**
- 3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.**

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;  
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

13. What is a Join Index?

- 1. Definition you should say first: A stored join or aggregation that the optimizer can use. Then spend the rest of the answer on architecture, a working example, and a production failure.**
- 2. A Join Index is a stored join or aggregation the optimizer can pick. It costs insert/update/delete maintenance.**
- 3. Worth it when a dashboard repeats the same join and you cannot change the app SQL. Recollect stats on the JI too.**

4. Working code / syntax (write this on Teradata SQL):

```
CREATE JOIN INDEX edw.ji_order_cust AS  
SELECT f.order_id, f.amt, c.segment  
FROM edw.fact_orders f  
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id  
PRIMARY INDEX (order_id);
```

5. A JI that is never chosen in EXPLAIN is dead weight. Drop it.

14. What is PPI?

- 1. Definition you should say first: Partitioned Primary Index: partitions rows inside each AMP. Then spend the rest of the answer on architecture, a working example, and a production failure.**
- 2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.**
- 3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.**

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE edw.fact_orders (  

```

JobKit - Teradata DBA interview questions

```
order_id BIGINT NOT NULL,  
order_dt DATE NOT NULL,  
amt DECIMAL(18,2)  
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);  
5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.
```

15. What is a NoPI table?

1. Definition you should say first: A table without a PI; used for fast staging loads. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.
3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE stg.orders_in (  
    order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)  
) NO PRIMARY INDEX;  
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE,  
    amt DECIMAL(18,2)  
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

16. What is a sparse map?

1. Definition you should say first: A map that uses a subset of AMPs for smaller tables. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Contiguous map = all AMPs. Sparse map = subset for small tables. Map change / reconfig redistributes objects after adding AMPs.
3. Do not put a 2-row calendar on 1000 AMPs without a sparse map if your version supports it. Plan reconfig with BAR and a freeze.

4. Working code / syntax (write this on Teradata SQL):

```
SHOW TABLE edw.dim_calendar;  
-- check MAP in newer Vantage  
SELECT DatabaseName, TableName, MapName FROM DBC.TablesV WHERE TableName = 'fact_orders';
```

5. Reconfig without a back-out window is a career incident.

17. What is a contiguous map?

1. Definition you should say first: The default map that uses all AMPs in the system. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Contiguous map = all AMPs. Sparse map = subset for small tables. Map change / reconfig redistributes objects after adding AMPs.
3. Do not put a 2-row calendar on 1000 AMPs without a sparse map if your version supports it. Plan reconfig with BAR and a freeze.

4. Working code / syntax (write this on Teradata SQL):

```
SHOW TABLE edw.dim_calendar;  
-- check MAP in newer Vantage  
SELECT DatabaseName, TableName, MapName FROM DBC.TablesV WHERE TableName = 'fact_orders';
```

5. Reconfig without a back-out window is a career incident.

18. What is hashing in Teradata?

1. Definition you should say first: The PI value is hashed to pick the AMP. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.
3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE stg.orders_in (  

```

JobKit - Teradata DBA interview questions

```
order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)
) NO PRIMARY INDEX;
CREATE MULTISET TABLE edw.fact_orders (
  order_id BIGINT NOT NULL,
  order_dt DATE,
  amt DECIMAL(18,2)
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

19. What is skew?

1. Definition you should say first: Uneven row or CPU distribution across AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Skew is uneven rows or CPU across AMPs. Low-cardinality PI (country, flag, null-heavy key) causes it.

3. Compare AMP row counts, DBQL, Viewpoint skew. Fix PI, add a composite PI, or stage and redistribute on a better key.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW(pi_col))) AS amp_n, COUNT(*) AS rows
FROM edw.fact_orders
GROUP BY 1
ORDER BY 2 DESC;
-- skew ratio = max/avg
```

5. If NULLs hash together, a nullable PI becomes one giant AMP. Coalesce or pick another key.

20. How do you find skewed tables?

1. Definition you should say first: Compare AMP row counts or use DBQL and Viewpoint. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Skew is uneven rows or CPU across AMPs. Low-cardinality PI (country, flag, null-heavy key) causes it.

3. Compare AMP row counts, DBQL, Viewpoint skew. Fix PI, add a composite PI, or stage and redistribute on a better key.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW(pi_col))) AS amp_n, COUNT(*) AS rows
FROM edw.fact_orders
GROUP BY 1
ORDER BY 2 DESC;
-- skew ratio = max/avg
```

5. If NULLs hash together, a nullable PI becomes one giant AMP. Coalesce or pick another key.

21. What is a hot AMP?

1. Definition you should say first: An AMP doing more work than others, often from skew. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew). Down AMP triggers fallback/cliue recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.

3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;
-- Viewpoint: AMP usage / skew
```

5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

22. What is spool space?

1. Definition you should say first: Temporary space for intermediate query results. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.

3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

23. What is perm space?

JobKit - Teradata DBA interview questions

1. Definition you should say first: Permanent table storage quota. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.
3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

24. What is temp space?

1. Definition you should say first: Space for global temporary tables. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.
3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

25. User vs database in Teradata?

1. Definition you should say first: A user can log on; a database is a space container. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.
3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

26. What is a profile in Teradata?

1. Definition you should say first: A set of account, spool, and password settings for users. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.
3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
GRANT role_read_edw TO u_jaya;
REVOKE role_read_edw FROM u_ex;
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

27. What is a role?

1. Definition you should say first: A bundle of rights you grant to users. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.
3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
GRANT role_read_edw TO u_jaya;
```

JobKit - Teradata DBA interview questions

```
REVOKE role_read_edw FROM u_ex;  
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

28. What is GRANT vs role?

1. Definition you should say first: GRANT is a privilege; a role groups many privileges. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.

3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;  
GRANT role_read_edw TO u_jaya;  
REVOKE role_read_edw FROM u_ex;  
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

29. What is a macro?

1. Definition you should say first: A stored SQL snippet executed with EXEC. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (  
    SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE  
) WITH DATA  
PRIMARY INDEX (order_id)  
ON COMMIT PRESERVE ROWS;  
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')  
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

30. What is a stored procedure?

1. Definition you should say first: Procedural SQL with logic, not just a macro. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (  
    SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE  
) WITH DATA  
PRIMARY INDEX (order_id)  
ON COMMIT PRESERVE ROWS;  
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')  
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

31. What is a volatile table?

1. Definition you should say first: Session-scoped table that disappears on logoff. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (  
    SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE
```

JobKit - Teradata DBA interview questions

```
) WITH DATA
PRIMARY INDEX (order_id)
ON COMMIT PRESERVE ROWS;
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

32. Global temporary vs volatile?

1. Definition you should say first: GTT definition is permanent; data is session-scoped. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (
  SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE
) WITH DATA
PRIMARY INDEX (order_id)
ON COMMIT PRESERVE ROWS;
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

33. What is a derived table?

1. Definition you should say first: An inline subquery in the FROM clause. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (
  SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE
) WITH DATA
PRIMARY INDEX (order_id)
ON COMMIT PRESERVE ROWS;
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

34. What is a CTE?

1. Definition you should say first: WITH clause query that names a result set. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.

3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (
  SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE
) WITH DATA
PRIMARY INDEX (order_id)
ON COMMIT PRESERVE ROWS;
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

35. What is COLLECT STATISTICS?

1. Definition you should say first: Gather data demographics so the optimizer can plan well. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.
3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.

4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);  
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);  
HELP STATISTICS ON edw.fact_orders;
```

5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

36. When do you collect stats?

1. Definition you should say first: After large loads, skew, or bad plans; not blindly every hour. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.

3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.

4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);  
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);  
HELP STATISTICS ON edw.fact_orders;
```

5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

37. What is SUMMARY stats?

1. Definition you should say first: A cheaper stats option on some releases for quick demographics. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.

3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.

4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);  
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);  
HELP STATISTICS ON edw.fact_orders;
```

5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

38. What is sampled stats?

1. Definition you should say first: Stats on a sample; faster but less accurate. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.

3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.

4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);  
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);  
HELP STATISTICS ON edw.fact_orders;
```

5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

39. How do you see a plan?

1. Definition you should say first: EXPLAIN SELECT ... Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN  
SELECT f.order_id, c.segment  
FROM edw.fact_orders f  
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id  
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

JobKit - Teradata DBA interview questions

5. SET tables with duplicate checks plus a product join will run until spool dies.

40. What is a confidence level in EXPLAIN?

1. Definition you should say first: How sure the optimizer is about row estimates. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

41. What is a product join?

1. Definition you should say first: Every row compared to every row; often a warning sign. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

42. What is a merge join?

1. Definition you should say first: Join on sorted/hashed keys; usually cheaper. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

43. What is a hash join?

1. Definition you should say first: Build a hash table of the smaller set and probe it. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

44. What is a nested join?

1. Definition you should say first: Index-driven lookup join. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

45. What is redistribution?

1. Definition you should say first: Moving rows across AMPs so join keys co-locate. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

46. What is duplication of a table in a join?

1. Definition you should say first: Copying a small table to all AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

47. What is a residual condition?

1. Definition you should say first: A filter applied after the join or retrieve step. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

48. What is a lock in Teradata?

1. Definition you should say first: Access, read, write, or exclusive lock on a table or row hash. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

JobKit - Teradata DBA interview questions

LOCKING ROW FOR ACCESS

```
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

49. What is a deadlock?

1. Definition you should say first: Two sessions waiting on each other; one is aborted. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

LOCKING ROW FOR ACCESS

```
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

50. How do you reduce lock contention?

1. Definition you should say first: Shorter transactions, nightly batch windows, and row-hash locks. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

LOCKING ROW FOR ACCESS

```
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

51. What is a transaction in Teradata?

1. Definition you should say first: BT/ET or ANSI transaction around SQL. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. BTET (Teradata mode) auto-commits each request unless BT/ET. ANSI uses BEGIN/COMMIT. A request is one SQL; a transaction can be many.

3. ETL scripts should be explicit: BT; steps; ET; with error handling in BTEQ.

4. Working code / syntax (write this on Teradata SQL):

.SET SESSION TRANSACTION BTET

BT;

DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-01';

INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;

ET;

5. A failed INSERT after DELETE without a transaction leaves an empty partition. That is a Sev1.

52. BTET vs ANSI mode?

1. Definition you should say first: BTET is Teradata transaction mode; ANSI is commit-oriented. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. BTET (Teradata mode) auto-commits each request unless BT/ET. ANSI uses BEGIN/COMMIT. A request is one SQL; a transaction can be many.

3. ETL scripts should be explicit: BT; steps; ET; with error handling in BTEQ.

4. Working code / syntax (write this on Teradata SQL):

.SET SESSION TRANSACTION BTET

BT;

DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-01';

INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;

ET;

5. A failed INSERT after DELETE without a transaction leaves an empty partition. That is a Sev1.

53. What is FastLoad?

1. Definition you should say first: Utility to load empty tables quickly with two phases. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

54. FastLoad limitation?

1. Definition you should say first: Target must be empty; no SI during load in classic use. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

55. What is MultiLoad?

1. Definition you should say first: Utility for insert/update/delete on populated tables. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

56. What is TPump?

JobKit - Teradata DBA interview questions

1. Definition you should say first: Low-volume continuous load with row-hash locks. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

57. What is FastExport?

1. Definition you should say first: High-volume export from Teradata. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

58. What is BTEQ?

1. Definition you should say first: CLI for SQL, scripts, and simple imports. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

59. What is TPT?

1. Definition you should say first: Teradata Parallel Transporter: modern load/export framework. Then spend the rest of the

JobKit - Teradata DBA interview questions

answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

60. TPT vs legacy utilities?

1. Definition you should say first: TPT can replace FastLoad/MultiLoad/FastExport with operators. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

61. What is a load slot?

1. Definition you should say first: A limit on concurrent load jobs. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

62. What is a checkpoint in MultiLoad?

1. Definition you should say first: A restart point after a failure. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

63. What is an error table?

1. Definition you should say first: Where utilities land rejected rows. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

64. What is a UV table?

1. Definition you should say first: Uniqueness violation table in MultiLoad. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

65. What is Viewpoint?

1. Definition you should say first: Web console for health, queries, and alerts. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters,

JobKit - Teradata DBA interview questions

exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

66. What is QueryGrid?

1. Definition you should say first: Fabric to query Hadoop, Oracle, and other systems from Teradata. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

67. What is TASM?

1. Definition you should say first: Teradata Active System Management: workload rules and throttles. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

68. What is a workload in TASM?

1. Definition you should say first: A class of queries with priority and limits. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

69. What is a throttle?

1. Definition you should say first: A cap on concurrency for a workload. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

JobKit - Teradata DBA interview questions

70. What is a filter in TASM?

1. Definition you should say first: A rule that rejects or delays some queries. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.
4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```
5. Everything in WD-Default means TASM is not actually implemented.

71. What is exception handling in TASM?

1. Definition you should say first: Move or abort queries that exceed CPU or skew. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.
4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```
5. Everything in WD-Default means TASM is not actually implemented.

72. What is WD-Default?

1. Definition you should say first: The default workload when nothing else matches. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.
4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```
5. Everything in WD-Default means TASM is not actually implemented.

73. What is a utility session?

1. Definition you should say first: Sessions used by load/export utilities. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.
4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```
5. Everything in WD-Default means TASM is not actually implemented.

74. What is BAR?

1. Definition you should say first: Backup and restore, often with DSA or NetBackup. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
```

JobKit - Teradata DBA interview questions

```
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

75. What is DSA?

1. Definition you should say first: Data Stream Architecture for backups. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

76. Full vs incremental backup?

1. Definition you should say first: Full copies everything; incremental copies changes. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

77. How do you restore one table?

1. Definition you should say first: Use BAR object-level restore if the backup supports it. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

78. What is a dump vs backup?

1. Definition you should say first: Dump often means a system dump for support; backup is data protection. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

79. What is PDBA vs production DBA?

1. Definition you should say first: PDBA may own platform; production DBA owns day-to-day jobs and users. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

80. What is space reclamation?

1. Definition you should say first: Freeing perm after deletes; packing and fallback copies matter. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.
3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

81. What is a cylinder?

1. Definition you should say first: A unit of disk allocation on an AMP. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

82. What is a table header?

1. Definition you should say first: Metadata stored with the table on each AMP. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

83. What is a journal?

1. Definition you should say first: Before/after image used for recovery on some tables. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

84. Permanent journal vs transient?

1. Definition you should say first: Permanent is extra protection; transient is for transactions. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDDBA.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```
5. Running checktable on a live 100 TB fact without a window can be an outage.

85. What is TVS?

1. Definition you should say first: Teradata Virtual Storage: storage virtualization layer. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDDBA.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```
5. Running checktable on a live 100 TB fact without a window can be an outage.

86. What is FSYS?

1. Definition you should say first: File system layer under AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PDDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.
3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDDBA.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```
5. Running checktable on a live 100 TB fact without a window can be an outage.

87. What is a node crash recovery?

1. Definition you should say first: Vprocs migrate; cliques and HSN determine downtime. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```
5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

88. What is a database restart?

1. Definition you should say first: TPA reset: sessions drop; need to reconnect. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```
5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

89. What is TPA?

1. Definition you should say first: Teradata parallel database process stack. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

90. What is PDE?

1. Definition you should say first: Parallel Database Extensions: OS layer for vprocs. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

91. What is a gateway?

1. Definition you should say first: Connection manager for client sessions. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.
4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

92. What is a session?

1. Definition you should say first: A logged-on PE connection from a client. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

93. Default session mode?

1. Definition you should say first: Depends on client; know if you are ANSI or Teradata mode. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

JobKit - Teradata DBA interview questions

94. What is a request vs a transaction?

1. Definition you should say first: A request is one SQL; a transaction may be many requests. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BTET (Teradata mode) auto-commits each request unless BT/ET. ANSI uses BEGIN/COMMIT. A request is one SQL; a transaction can be many.
3. ETL scripts should be explicit: BT; steps; ET; with error handling in BTEQ.
4. Working code / syntax (write this on Teradata SQL):

```
.SET SESSION TRANSACTION BTET
BT;
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-01';
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;
ET;
```

5. A failed INSERT after DELETE without a transaction leaves an empty partition. That is a Sev1.

95. How do you kill a bad query?

1. Definition you should say first: Abort in Viewpoint or with a console abort after checking impact. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

96. How do you find who filled the spool?

1. Definition you should say first: DBQL plus user spool limits. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.
3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

97. What is DBQL?

1. Definition you should say first: Database Query Log: history of queries and steps. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

98. What is ResUsage?

1. Definition you should say first: Resource usage macros for CPU, disk, and BYNET. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
```

JobKit - Teradata DBA interview questions

ABORT SESSION 1, 1042;

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

99. What is a heart beat in Viewpoint?

1. Definition you should say first: System health pulse; red means investigate now. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

100. What is a blocked session?

1. Definition you should say first: Waiting on a lock held by another session. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS  
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

101. How do you see blockers?

1. Definition you should say first: Lock viewer in Viewpoint or lock query views. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS  
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

102. What is a deadlock timeout?

1. Definition you should say first: How long Teradata waits before aborting a deadlock victim. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS  
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;  
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo  
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

103. What is a PPI advantage?

1. Definition you should say first: Partition elimination so fewer cylinders are read. Then spend the rest of the answer on

JobKit - Teradata DBA interview questions

architecture, a working example, and a production failure.

2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.

3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE NOT NULL,  
    amt DECIMAL(18,2)  
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);
```

5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.

104. What is a character PPI?

1. Definition you should say first: Partitioning on a character column; less common than date PPI. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.

3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE NOT NULL,  
    amt DECIMAL(18,2)  
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);
```

5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.

105. What is multilevel PPI?

1. Definition you should say first: More than one partition expression. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.

3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE NOT NULL,  
    amt DECIMAL(18,2)  
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);
```

5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.

106. What is a column partition (CP)?

1. Definition you should say first: Columnar storage for analytic tables on some releases. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.

3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE NOT NULL,  
    amt DECIMAL(18,2)
```

JobKit - Teradata DBA interview questions

```
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);  
5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.
```

107. What is Hybrid Row/Column?

1. Definition you should say first: Mix of row and column partitions. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PPI partitions rows inside each AMP (often by date) so scans prune. Character PPI is possible but date/int is cleaner. CP / hybrid are columnar options for analytic compression.
3. Partition facts by order_dt RANGE_N. Collect stats on the partition column. Avoid functions on the partition column in WHERE.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISSET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE NOT NULL,  
    amt DECIMAL(18,2)  
)  
PRIMARY INDEX (order_id)  
PARTITION BY RANGE_N(order_dt BETWEEN DATE '2018-01-01' AND DATE '2028-12-31' EACH INTERVAL '1' MONTH);  
5. WHERE EXTRACT(YEAR FROM order_dt)=2026 kills pruning. Use a date range.
```

108. What is compression in Teradata?

1. Definition you should say first: MVC, UDT, algorithmic, or block-level depending on release. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.
3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

109. What is MVC?

1. Definition you should say first: Multi-Value Compression: dictionary compression on a column. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.
3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

110. What is a UDT?

1. Definition you should say first: User-defined type. Then spend the rest of the answer on architecture, a working example, and a production failure.
 2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.
 3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.
- #### 4. Working code / syntax (write this on Teradata SQL):

JobKit - Teradata DBA interview questions

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

111. What is a UDF?

1. Definition you should say first: User-defined function. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.

3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

112. What is a table operator?

1. Definition you should say first: SQL-HMR style operator for custom processing. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.

3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

113. What is Native Object Store (NOS)?

1. Definition you should say first: Read/write object storage (S3-style) from Teradata Vantage. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. MVC = multi-value compression on repeating values. UDT/UDF extend types/functions. Table operators and NOS read object store (S3-style) from SQL.

3. Compress low-cardinality columns. NOS for lake files; Teradata remains the governed EDW.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE TABLE edw.dim_status (  
    status_cd CHAR(3) COMPRESS ('NEW', 'OPN', 'CLS'),  
    status_desc VARCHAR(40)  
) PRIMARY INDEX (status_cd);  
SELECT $PATH, col1 FROM (  
    LOCATION '/s3/bucket/path'  
) AS nos_d;
```

5. Compressing a near-unique column wastes CPU for no disk win.

114. What is Vantage?

1. Definition you should say first: Teradata's analytics platform branding for newer releases. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata is a shared-nothing MPP analytic database. Vantage is the current platform branding (SQL plus analytics). Each AMP owns a slice of every table.

JobKit - Teradata DBA interview questions

3. Draw PE -> BYNET -> AMPs. Mention PI, stats, TASM, and BAR in the first minute.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.DiskSpaceV WHERE DatabaseName = 'EDW';  
HELP SESSION;  
SELECT HASHAMP() AS amps;
```

5. Calling it 'just like Oracle but faster' without AMP/PI/skew will not survive a DBA round.

115. What is QueryBand?

1. Definition you should say first: Name/value tags on a session or transaction for workload and audit. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

116. What is an account string?

1. Definition you should say first: Account used for chargeback and TASM mapping. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

117. What is a performance group?

1. Definition you should say first: Older priority scheme; TASM workloads replaced much of this. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

118. What is a resource partition?

1. Definition you should say first: A slice of CPU reserved for a set of workloads. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.

3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

119. What is tactical vs decision support?

JobKit - Teradata DBA interview questions

1. Definition you should say first: Tactical is short; DSS is heavy reporting. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

120. How do you protect tactical queries?

1. Definition you should say first: TASM: high priority, low throttle, and exception on long CPU. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

121. What is a delayed query?

1. Definition you should say first: Queued by a throttle until a slot is free. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

122. What is a reject in TASM?

1. Definition you should say first: The query is not allowed to run. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

123. What is an object lock vs row hash lock?

1. Definition you should say first: Object is table-level; row hash is finer. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.
3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
```

JobKit - Teradata DBA interview questions

```
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

124. What is a dummy row lock?

1. Definition you should say first: Locking a hash with no row; still blocks that hash. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

125. How do you load a large fact table daily?

1. Definition you should say first: Stage NoPI or FastLoad, then INSERT-SELECT with collect stats. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata is shared-nothing MPP: PE parses, BYNET ships, AMPs own their rows and spool.

3. Always name PI, stats, skew, spool, and the utility you would use. Then EXPLAIN before you run in prod.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT * FROM edw.fact_orders
WHERE order_dt = DATE '2026-09-01';
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id, order_dt);
```

5. A product join plus missing stats plus a bad NUPI will light up one hot AMP and fill spool.

126. What is an INSERT-SELECT?

1. Definition you should say first: Load from one table into another set-based. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

127. What is MERGE INTO?

1. Definition you should say first: Upsert based on a match condition. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
```

JobKit - Teradata DBA interview questions

```
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

128. What is UPDATE from a join?

1. Definition you should say first: Set-based update using another table. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

129. What is a SET vs MULTISET table?

1. Definition you should say first: SET rejects duplicate rows; MULTISET allows them. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SET tables reject duplicate rows (expensive duplicate row check). MULTISET allows duplicates and is the default for most ETL facts.

3. Use SET only when uniqueness must be enforced by the table and volume is modest. Prefer UPI or USI + MULTISET for facts.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE SET TABLE edw.dim_code (
    code CHAR(4) NOT NULL,
    descr VARCHAR(40)
) UNIQUE PRIMARY INDEX (code);
CREATE MULTISET TABLE edw.fact_click (click_id BIGINT, url VARCHAR(200)) PRIMARY INDEX (click_id);
```

5. A SET fact with no UPI will spend CPU comparing rows on every insert.

130. Duplicate row check cost?

1. Definition you should say first: SET tables can be expensive on NUPI during inserts. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SET tables reject duplicate rows (expensive duplicate row check). MULTISET allows duplicates and is the default for most ETL facts.

3. Use SET only when uniqueness must be enforced by the table and volume is modest. Prefer UPI or USI + MULTISET for facts.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE SET TABLE edw.dim_code (
    code CHAR(4) NOT NULL,
    descr VARCHAR(40)
) UNIQUE PRIMARY INDEX (code);
CREATE MULTISET TABLE edw.fact_click (click_id BIGINT, url VARCHAR(200)) PRIMARY INDEX (click_id);
```

5. A SET fact with no UPI will spend CPU comparing rows on every insert.

131. What is a queue table?

JobKit - Teradata DBA interview questions

1. Definition you should say first: A table with consume semantics for messaging-style pulls. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.
3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (  
    SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE  
) WITH DATA  
PRIMARY INDEX (order_id)  
ON COMMIT PRESERVE ROWS;  
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')  
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

132. What is an identity column?

1. Definition you should say first: System-generated numbers; be careful with distribution. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Macro = parameterized SQL. SP = procedural. Volatile = session table. GTT = permanent def, session data. Derived = FROM (SELECT...). CTE = WITH. Queue tables support FIFO. Identity is a generated number, not a PI by default thinking.
3. Use volatile for intermediate in ETL scripts. Do not hide 2000-line business logic only in SPs if the team cannot EXPLAIN it.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VOLATILE TABLE vt_today AS (  
    SELECT * FROM edw.fact_orders WHERE order_dt = CURRENT_DATE  
) WITH DATA  
PRIMARY INDEX (order_id)  
ON COMMIT PRESERVE ROWS;  
WITH c AS (SELECT cust_id FROM edw.dim_customer WHERE segment = 'GOLD')  
SELECT * FROM edw.fact_orders f JOIN c ON c.cust_id = f.cust_id;
```

5. Volatile tables disappear on logoff - BTEQ scripts that reconnect lose them.

133. What is a date PI problem?

1. Definition you should say first: Low cardinality dates skew if used as PI. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.
3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE stg.orders_in (  
    order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)  
) NO PRIMARY INDEX;  
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE,  
    amt DECIMAL(18,2)  
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

134. Good PI for a fact table?

1. Definition you should say first: A high-cardinality join key, often a customer or transaction id. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. PI is the distribution key: HASHROW -> AMP. UPI unique, NUPI allows dupes. NoPI is unhashed staging. Date-only PI often skews. Fact PI is usually a high-cardinality id (order_id), not a status flag.
3. Choose PI for join geography and even AMP counts, not for 'it looks like a PK'. Stage on NoPI, then INSERT-SELECT to the PI table.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE MULTISET TABLE stg.orders_in (  
    order_id BIGINT, order_dt DATE, amt DECIMAL(18,2)
```

JobKit - Teradata DBA interview questions

```
) NO PRIMARY INDEX;  
CREATE MULTISET TABLE edw.fact_orders (  
    order_id BIGINT NOT NULL,  
    order_dt DATE,  
    amt DECIMAL(18,2)  
) PRIMARY INDEX (order_id);
```

5. Changing PI means a full rebuild. Get it right on day one of the model.

135. What is referential integrity in Teradata?

1. Definition you should say first: Optional; soft RI is often used for optimizer hints. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.
3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS  
SELECT f.*, c.segment  
FROM edw.fact_orders f  
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;  
-- SCD2 row:  
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

136. What is a batch vs tactical window?

1. Definition you should say first: Batch at night with loads; tactical during business hours. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Teradata is shared-nothing MPP: PE parses, BYNET ships, AMPs own their rows and spool.
3. Always name PI, stats, skew, spool, and the utility you would use. Then EXPLAIN before you run in prod.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN  
SELECT * FROM edw.fact_orders  
WHERE order_dt = DATE '2026-09-01';  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id, order_dt);
```

5. A product join plus missing stats plus a bad NUPI will light up one hot AMP and fill spool.

137. How do you handle a failed FastLoad?

1. Definition you should say first: Drop or empty the target, fix the file, restart; check error tables. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.
3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */  
-- DEFINE OPERATOR DDL / LOAD  
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)  
INSERT INTO edw.fact_orders  
SELECT * FROM stg.orders_in;  
MERGE INTO edw.dim_customer t  
USING stg.customer s ON t.cust_id = s.cust_id  
WHEN MATCHED THEN UPDATE SET name = s.name  
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

138. What is a two-phase FastLoad?

1. Definition you should say first: Phase 1 loading to AMPs; phase 2 sorting into the table. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework.

JobKit - Teradata DBA interview questions

INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

139. Can FastLoad load a PI with SI?

1. Definition you should say first: Classic FastLoad wants no secondary indexes on the target. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

140. What is a MiniBatch?

1. Definition you should say first: A pattern of small MultiLoad or TPT jobs. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

141. What is streaming ingest?

1. Definition you should say first: TPump or TPT Stream for near-real-time. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. FastLoad = empty table, two phases, no SI, max one PI. MultiLoad = maintained tables, upsert, checkpoints, error/UV tables. TPump = low volume trickle. FastExport = out. BTEQ = script/SQL. TPT = modern operator framework. INSERT-SELECT and MERGE are SQL loads.

JobKit - Teradata DBA interview questions

3. Nightly facts: TPT Load or FastLoad to NoPI/stg then INSERT-SELECT. Failed FastLoad: drop/release lock, clean err tables, restart from checkpoint rules.

4. Working code / syntax (write this on Teradata SQL):

```
/* TPT Load idea */
-- DEFINE OPERATOR DDL / LOAD
-- FastLoad limits: empty table, no SI, no PPI in old rules (version-specific)
INSERT INTO edw.fact_orders
SELECT * FROM stg.orders_in;
MERGE INTO edw.dim_customer t
USING stg.customer s ON t.cust_id = s.cust_id
WHEN MATCHED THEN UPDATE SET name = s.name
WHEN NOT MATCHED THEN INSERT (cust_id, name) VALUES (s.cust_id, s.name);
```

5. Leaving a FastLoad lock on a table blocks the morning. Always script the release/cleanup path.

142. How do you grant access to a new analyst?

1. Definition you should say first: Role-based: create user, grant role, set spool and profile. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.

3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
GRANT role_read_edw TO u_jaya;
REVOKE role_read_edw FROM u_ex;
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

143. How do you revoke a leaver?

1. Definition you should say first: Lock the user, revoke roles, drop or keep for audit per policy. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.

3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
GRANT role_read_edw TO u_jaya;
REVOKE role_read_edw FROM u_ex;
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

144. What is a password lockout?

1. Definition you should say first: Profile setting after failed logons. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.

3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
GRANT role_read_edw TO u_jaya;
REVOKE role_read_edw FROM u_ex;
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

145. What is LDAP vs TD database auth?

1. Definition you should say first: Enterprise directory vs Teradata passwords. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.

3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;
```

JobKit - Teradata DBA interview questions

```
GRANT role_read_edw TO u_jaya;  
REVOKE role_read_edw FROM u_ex;  
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

146. What is SSO to Viewpoint?

1. Definition you should say first: Viewpoint can use corporate identity; still map TD users. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. A profile holds password/spool defaults. A role is a bundle of rights. GRANT to roles, not 200 users. LDAP/SSO for people; TD passwords for service ids.
3. Joiner: create user, grant role_analyst, set profile. Leaver: revoke role, lock, then drop after HR date. Never share DBC.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE USER u_jaya FROM edw_users AS PASSWORD = *** PERM = 0 SPOOL = 5e9 PROFILE = pr_analyst;  
GRANT role_read_edw TO u_jaya;  
REVOKE role_read_edw FROM u_ex;  
MODIFY USER u_ex AS PASSWORD = *** LOCKED = U;
```

5. Personal copies of prod tables in a sandbox with PII is an audit finding.

147. What is a zone in Viewpoint?

1. Definition you should say first: A filtered view of systems for a team. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

148. What is a portlet?

1. Definition you should say first: A Viewpoint dashboard widget. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Viewpoint is the ops UI (portlets, zones, heartbeat). TASM is workload management: workloads, throttles, filters, exceptions. Tactical = short SLA; DSS = heavy. QueryBand tags sessions. QueryGrid talks to other platforms.
3. Put tactical in a high-priority partition with a throttle on DSS. Reject or delay runaways. Set QueryBand in ETL for chargeback.

4. Working code / syntax (write this on Teradata SQL):

```
SET QUERY_BAND = 'App=ETL;Job=FACT_ORDERS;Team=EDW;' FOR SESSION;  
-- TASM: WD_Tactical timeout 10s, throttle DSS to N queries  
-- Viewpoint alert: heartbeat down, blocked sessions > 20, spool > 80%
```

5. Everything in WD-Default means TASM is not actually implemented.

149. What is a system heartbeat alert?

1. Definition you should say first: Pager-worthy if the database is down. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Teradata is shared-nothing MPP: PE parses, BYNET ships, AMPs own their rows and spool.
3. Always name PI, stats, skew, spool, and the utility you would use. Then EXPLAIN before you run in prod.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN  
SELECT * FROM edw.fact_orders  
WHERE order_dt = DATE '2026-09-01';  
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id, order_dt);
```

5. A product join plus missing stats plus a bad NUPI will light up one hot AMP and fill spool.

150. How do you plan an upgrade?

1. Definition you should say first: Read the release, test BAR, check views, schedule an outage window. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Upgrades need compat checks, BAR, a freeze, and a back-out. Stall in Viewpoint is wait time, not always a down system.
3. Read release notes for TPT/TASM changes. Run CHECKTABLE only per vendor guidance.

JobKit - Teradata DBA interview questions

4. Working code / syntax (write this on Teradata SQL):

```
-- pre-upgrade
-- 1 BAR full 2 capture TASM 3 freeze 4 upgrade 5 validate Viewpoint heartbeat 6 smoke FastLoad + SELECT
HASHAMP()
```

5. Skipping a backup because 'it is only a patch' is the story in half of DBA incident talks.

151. What is a compat view?

1. Definition you should say first: Older dictionary views kept for scripts. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. DBC is the dictionary. SYSLIB has functions. *V views are the supported interface. Always query views, not legacy base tables, unless PDBA asks.

3. Size from TableSizeV / DiskSpaceV. Space from DiskSpaceV CurrentPerm vs MaxPerm.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT DatabaseName, TableName, SUM(CurrentPerm) AS bytes
FROM DBC.TableSizeV
WHERE DatabaseName = 'EDW'
GROUP BY 1, 2
ORDER BY 3 DESC;
SELECT TableName, TableKind FROM DBC.TablesV WHERE DatabaseName = 'EDW';
```

5. Updating DBC base tables by hand is how you corrupt a system.

152. DBC vs SYSLIB?

1. Definition you should say first: DBC is the data dictionary; SYSLIB holds functions. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. DBC is the dictionary. SYSLIB has functions. *V views are the supported interface. Always query views, not legacy base tables, unless PDBA asks.

3. Size from TableSizeV / DiskSpaceV. Space from DiskSpaceV CurrentPerm vs MaxPerm.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT DatabaseName, TableName, SUM(CurrentPerm) AS bytes
FROM DBC.TableSizeV
WHERE DatabaseName = 'EDW'
GROUP BY 1, 2
ORDER BY 3 DESC;
SELECT TableName, TableKind FROM DBC.TablesV WHERE DatabaseName = 'EDW';
```

5. Updating DBC base tables by hand is how you corrupt a system.

153. What is DBC.TablesV?

1. Definition you should say first: Dictionary view of tables. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. DBC is the dictionary. SYSLIB has functions. *V views are the supported interface. Always query views, not legacy base tables, unless PDBA asks.

3. Size from TableSizeV / DiskSpaceV. Space from DiskSpaceV CurrentPerm vs MaxPerm.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT DatabaseName, TableName, SUM(CurrentPerm) AS bytes
FROM DBC.TableSizeV
WHERE DatabaseName = 'EDW'
GROUP BY 1, 2
ORDER BY 3 DESC;
SELECT TableName, TableKind FROM DBC.TablesV WHERE DatabaseName = 'EDW';
```

5. Updating DBC base tables by hand is how you corrupt a system.

154. How do you find table size?

1. Definition you should say first: SUM of currentperm from DBC.TableSizeV grouped by table. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. DBC is the dictionary. SYSLIB has functions. *V views are the supported interface. Always query views, not legacy base tables, unless PDBA asks.

3. Size from TableSizeV / DiskSpaceV. Space from DiskSpaceV CurrentPerm vs MaxPerm.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT DatabaseName, TableName, SUM(CurrentPerm) AS bytes
FROM DBC.TableSizeV
WHERE DatabaseName = 'EDW'
```

JobKit - Teradata DBA interview questions

```
GROUP BY 1, 2
ORDER BY 3 DESC;
SELECT TableName, TableKind FROM DBC.TablesV WHERE DatabaseName = 'EDW';
```

5. Updating DBC base tables by hand is how you corrupt a system.

155. How do you find user space?

1. Definition you should say first: DBC.DiskSpaceV or Viewpoint space portlet. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. DBC is the dictionary. SYSLIB has functions. *V views are the supported interface. Always query views, not legacy base tables, unless PDBA asks.

3. Size from TableSizeV / DiskSpaceV. Space from DiskSpaceV CurrentPerm vs MaxPerm.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT DatabaseName, TableName, SUM(CurrentPerm) AS bytes
FROM DBC.TableSizeV
WHERE DatabaseName = 'EDW'
GROUP BY 1, 2
ORDER BY 3 DESC;
SELECT TableName, TableKind FROM DBC.TablesV WHERE DatabaseName = 'EDW';
```

5. Updating DBC base tables by hand is how you corrupt a system.

156. What is currentperm vs peakperm?

1. Definition you should say first: Current used vs high-water used. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Perm = table bytes. Spool = query scratch. Temp = GTT data. Users log on; databases hold objects. PeakPerm is high-water. Reclamation frees deleted space over time.

3. When spool exceeds, find the query in DBQL, EXPLAIN, add stats, or raise spool with a ticket - do not silently grant unlimited.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT UserName, PeakSpool FROM DBC.DiskSpaceV WHERE PeakSpool > 0 ORDER BY 2 DESC;
SELECT DatabaseName, SUM(CurrentPerm) FROM DBC.TableSizeV GROUP BY 1;
-- abort: ABORT SESSION <host>, <session>
```

5. Granting huge spool to hide a product join is not tuning.

157. What is a stall?

1. Definition you should say first: AMPs waiting on I/O or each other; check ResUsage. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.

3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```

5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

158. What is a bottleneck on BYNET?

1. Definition you should say first: Heavy redistribution; redesign PI or join. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. BYNET is the interconnect for messages, row redistribution, and duplication between PEs and AMPs.

3. Product joins and redistributing a huge uncompressed table flood BYNET. Check ResUsage and Viewpoint fabric metrics.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT *
FROM edw.fact_orders f
JOIN edw.dim_customer d ON d.region = f.region;
-- look for: redistributing / duplicating all rows of
```

5. If you only add CPU and ignore join geography, BYNET will still stall.

159. What is a CPU-bound query?

1. Definition you should say first: High AMP CPU; check product joins and functions on columns. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. Sessions enter via gateway. Default mode is often BTET (site-specific). DBQL logs queries. ResUsage is resource telemetry. Stall = waiting (locks, I/O). CPU-bound = AMP CPU high with little I/O wait.
3. Kill with abort session after you capture EXPLAIN/DBQL. Do not kill a BAR or a FastLoad without the utility owner.
4. Working code / syntax (write this on Teradata SQL):

```
SELECT SessionNo, UserName, StatementText FROM DBC.QryLogV
WHERE StartTime > CURRENT_TIMESTAMP - INTERVAL '10' MINUTE;
ABORT SESSION 1, 1042;
```
5. Killing the wrong session during MERGE can leave a transaction open. Check SessionInfo first.

160. What is a missing stats symptom?

1. Definition you should say first: Wrong join type and huge spool. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.
3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.
4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);
HELP STATISTICS ON edw.fact_orders;
```
5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

161. How often to collect stats on a large fact?

1. Definition you should say first: After major loads; use THRESHOLD where available. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.
3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.
4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);
HELP STATISTICS ON edw.fact_orders;
```
5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

162. What is stats THRESHOLD?

1. Definition you should say first: Skip collect if data has not changed enough. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Stats are demographics for the optimizer. Collect after large loads or skew. Sampled/summary are cheaper. THRESHOLD skips collect if little change. Huge facts: collect on PI, partition, join columns - not every column daily.
3. Help stats, compare EXPLAIN estimates vs actual DBQL. Recollect when a job's plan flips.
4. Working code / syntax (write this on Teradata SQL):

```
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_id);
COLLECT STATISTICS ON edw.fact_orders COLUMN (order_dt);
COLLECT STATISTICS USING SAMPLE ON edw.stg_big COLUMN (load_dt);
HELP STATISTICS ON edw.fact_orders;
```
5. Collecting all columns hourly on a 20 TB fact is a worse outage than stale stats.

163. What is a copy job in BAR?

1. Definition you should say first: Copy backup data to another system. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```
5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

JobKit - Teradata DBA interview questions

164. What is a restore to a different system?

1. Definition you should say first: Needs compatible maps and enough AMPs or a remap strategy. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. BAR = backup and restore. DSA is the modern BAR. Full vs incremental. Dump is often a looser term; production uses catalogued BAR jobs. Restore one table for a bad load; restore to another system for DR tests.
3. Nightly incremental + weekly full. Test a table restore quarterly. Document the job name in the runbook.
4. Working code / syntax (write this on Teradata SQL):

```
-- DSA / BAR job (logical)
JOB: BAR_EDW_DAILY
  INCLUDE DATABASE edw
  INCREMENTAL
RESTORE TABLE edw.fact_orders FROM BAR_EDW_DAILY WHERE BACKUP_DATE = DATE '2026-09-02';
```

5. A backup that never had a restore drill is not a backup. Interviewers ask when you last restored.

165. What is a map change?

1. Definition you should say first: Redistribute tables after adding AMPs. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Contiguous map = all AMPs. Sparse map = subset for small tables. Map change / reconfig redistributes objects after adding AMPs.
3. Do not put a 2-row calendar on 1000 AMPs without a sparse map if your version supports it. Plan reconfig with BAR and a freeze.
4. Working code / syntax (write this on Teradata SQL):

```
SHOW TABLE edw.dim_calendar;
-- check MAP in newer Vantage
SELECT DatabaseName, TableName, MapName FROM DBC.TablesV WHERE TableName = 'fact_orders';
```

5. Reconfig without a back-out window is a career incident.

166. What is a reconfig?

1. Definition you should say first: System change such as adding nodes; planned outage. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. Contiguous map = all AMPs. Sparse map = subset for small tables. Map change / reconfig redistributes objects after adding AMPs.
3. Do not put a 2-row calendar on 1000 AMPs without a sparse map if your version supports it. Plan reconfig with BAR and a freeze.

4. Working code / syntax (write this on Teradata SQL):

```
SHOW TABLE edw.dim_calendar;
-- check MAP in newer Vantage
SELECT DatabaseName, TableName, MapName FROM DBC.TablesV WHERE TableName = 'fact_orders';
```

5. Reconfig without a back-out window is a career incident.

167. What is HSN vs clique failover?

1. Definition you should say first: HSN is a spare node; clique shares disks among nodes. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.
3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.

4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

168. What is a down AMP?

1. Definition you should say first: AMP not processing; fallback or HSN may cover. Then spend the rest of the answer on architecture, a working example, and a production failure.
2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew). Down AMP triggers fallback/clique recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.
3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;
```

JobKit - Teradata DBA interview questions

```
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;
-- Viewpoint: AMP usage / skew
```

5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

169. What is a catchup?

1. Definition you should say first: Recovering fallback copies after a component returns. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. vproc = virtual processor (AMP/PE). Clique = nodes sharing disks. HSN = spare node. PDE/TPA = OS/DB layers. Gateway accepts sessions. Catchup replays after AMP recovery.

3. On node loss, vprocs migrate in the clique or onto HSN. Users see a restart window. Check Viewpoint heartbeat.

4. Working code / syntax (write this on Teradata SQL):

```
-- After restart:
SELECT * FROM DBC.VprocStatus;
SELECT * FROM DBC.DatabaseStatus;
```

5. No fallback + down AMP + no clique spare = table inaccessible. Know your availability design.

170. What is a checktable?

1. Definition you should say first: Utility to check table integrity. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.

3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

171. What is a scandisk?

1. Definition you should say first: Lower-level disk integrity check. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.

3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

172. What is a perm cache?

1. Definition you should say first: Memory cache of perm data; watch hit rates. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PE = Parsing Engine: session handling, parse, optimize, dispatch steps to AMPs.

3. PE CPU spikes on bad dynamic SQL or too many short tactical requests. TASM protects mixed workloads.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT AccountName, UserName FROM DBC.SessionInfoV WHERE UserName = USER;
EXPLAIN SELECT COUNT(*) FROM edw.fact_orders;
```

5. People confuse PE with AMP. PE does not store the table; AMPs do.

173. What is FSG cache?

1. Definition you should say first: File segment cache for I/O. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement. Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.

3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

JobKit - Teradata DBA interview questions

5. Running checktable on a live 100 TB fact without a window can be an outage.

174. What is a TVS temperature?

1. Definition you should say first: Hot/cold data placement on mixed storage. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. PDDBA is platform DBA (OS, PDE, BAR). Cylinders/FSG/TVS are storage layers and temperature-aware placement.

Journals: transient for tx, permanent optional. Checktable/scandisk are integrity tools used with care.

3. You will not run scandisk on a hunch in prod. You will watch TVS temperature and FSG hit rate with the PDDBA.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT * FROM DBC.StorageSizeV SAMPLE 10;
-- FSG cache hit: ResUsage / Viewpoint memory portlet
-- Permanent journal: show table will list JOURNAL options
```

5. Running checktable on a live 100 TB fact without a window can be an outage.

175. What is a 1-AMP operation?

1. Definition you should say first: A query that hits a single AMP via unique PI access. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew).

Down AMP triggers fallback/cliquest recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.

3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;
-- Viewpoint: AMP usage / skew
```

5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

176. What is an all-AMP retrieve?

1. Definition you should say first: Every AMP reads its piece of the table. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew).

Down AMP triggers fallback/cliquest recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.

3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;
-- Viewpoint: AMP usage / skew
```

5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

177. What is a group AMP operation?

1. Definition you should say first: A subset of AMPs, for example sparse maps. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. AMP = Access Module Processor. It stores rows, executes steps, and uses spool. A hot AMP does extra work (skew).

Down AMP triggers fallback/cliquest recovery. 1-AMP ops hit one AMP; all-AMP retrieve hits every AMP.

3. Check AMP CPU and row counts when a query is slow. Hash the PI to see placement.

4. Working code / syntax (write this on Teradata SQL):

```
SELECT HASHAMP(HASHBUCKET(HASHROW('1001')))) AS amp_guess;
SELECT vproc, COUNT(*) FROM edw.fact_orders GROUP BY 1 ORDER BY 2 DESC;
-- Viewpoint: AMP usage / skew
```

5. A UPI on gender or date can put millions of rows on one AMP. That is the classic hot AMP.

178. What is a join index maintenance cost?

1. Definition you should say first: Extra work on every base-table change. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A Join Index is a stored join or aggregation the optimizer can pick. It costs insert/update/delete maintenance.

3. Worth it when a dashboard repeats the same join and you cannot change the app SQL. Recollect stats on the JI too.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE JOIN INDEX edw.ji_order_cust AS
SELECT f.order_id, f.amt, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
PRIMARY INDEX (order_id);
```

JobKit - Teradata DBA interview questions

5. A JI that is never chosen in EXPLAIN is dead weight. Drop it.

179. When is a JI worth it?

1. Definition you should say first: Stable, repeated joins that dominate CPU. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A Join Index is a stored join or aggregation the optimizer can pick. It costs insert/update/delete maintenance.

3. Worth it when a dashboard repeats the same join and you cannot change the app SQL. Recollect stats on the JI too.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE JOIN INDEX edw.ji_order_cust AS
SELECT f.order_id, f.amt, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
PRIMARY INDEX (order_id);
```

5. A JI that is never chosen in EXPLAIN is dead weight. Drop it.

180. What is an aggregate JI?

1. Definition you should say first: Pre-aggregated stored result. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A Join Index is a stored join or aggregation the optimizer can pick. It costs insert/update/delete maintenance.

3. Worth it when a dashboard repeats the same join and you cannot change the app SQL. Recollect stats on the JI too.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE JOIN INDEX edw.ji_order_cust AS
SELECT f.order_id, f.amt, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
PRIMARY INDEX (order_id);
```

5. A JI that is never chosen in EXPLAIN is dead weight. Drop it.

181. What is a single-table JI?

1. Definition you should say first: A stored projection or extra access path on one table. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. A Join Index is a stored join or aggregation the optimizer can pick. It costs insert/update/delete maintenance.

3. Worth it when a dashboard repeats the same join and you cannot change the app SQL. Recollect stats on the JI too.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE JOIN INDEX edw.ji_order_cust AS
SELECT f.order_id, f.amt, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
PRIMARY INDEX (order_id);
```

5. A JI that is never chosen in EXPLAIN is dead weight. Drop it.

182. What is a covering index?

1. Definition you should say first: An index that satisfies a query without the base table. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.

3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

183. What is a value-ordered NUSI?

1. Definition you should say first: NUSI stored in value order for range access. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.

3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

JobKit - Teradata DBA interview questions

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

184. What is a hash-ordered NUSI?

1. Definition you should say first: Default NUSI organization. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.

3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;  
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

185. What is a unique secondary index?

1. Definition you should say first: Enforces uniqueness and supports lookups. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.

3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;  
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

186. Cost of USI on load?

1. Definition you should say first: Load utilities may require dropping USI first. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. SI is an extra access path with a subtable. USI enforces uniqueness. NUSI can be value-ordered or hash-ordered. Covering means the SI answers the query without the base row. USI/NUSI slow loads.

3. Add SI only for a proven tactical access path. Drop unused SI after DBQL review.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE UNIQUE INDEX usi_cust_email (email) ON edw.dim_customer;  
CREATE INDEX nusi_order_dt (order_dt) ON edw.fact_orders;
```

5. A USI on a FastLoad target is illegal or forces a slower path. Know utility limits.

187. What is a soft RI?

1. Definition you should say first: Optimizer uses the relationship without enforcing it. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS  
SELECT f.*, c.segment  
FROM edw.fact_orders f  
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;  
-- SCD2 row:  
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

188. What is a batch error table?

1. Definition you should say first: Landing for failed rows in a batch. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
```

JobKit - Teradata DBA interview questions

```
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

189. How do you handle late-arriving dimensions?

1. Definition you should say first: Staging and MERGE; do not invent keys. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

190. What is slowly changing dimension?

1. Definition you should say first: Type 1 overwrite vs Type 2 history rows. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

191. How does a DBA support SCD2?

1. Definition you should say first: PI, PPI on dates, and space for history. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

192. What is a surrogate key?

1. Definition you should say first: Internal integer key; pick a PI that does not skew. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then

JobKit - Teradata DBA interview questions

an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

193. What is a natural key?

1. Definition you should say first: Business key such as account number. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

194. What is 3NF vs star schema?

1. Definition you should say first: 3NF is normalized EDW; star is dimensional mart. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

195. Where does Teradata sit in a lakehouse?

1. Definition you should say first: Often the governed EDW; NOS can query the lake. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

196. What is data lineage?

1. Definition you should say first: Where a column came from; DBA supports via views and docs. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

197. What is a view vs a table?

1. Definition you should say first: View is stored SQL; table is stored rows. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

198. What is a recursive view?

1. Definition you should say first: View that refers to itself; use carefully. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Teradata RI can be hard or soft. SCD2 keeps history with surrogate keys. Star vs 3NF is mart vs EDW. NOS/lake sits beside the EDW. Views store SQL, not rows.

3. DBA supports SCD2 with space, PI on surrogate, PPI on dates, and stats. Late dimensions get a default unknown key then an update.

4. Working code / syntax (write this on Teradata SQL):

```
CREATE VIEW edw.v_orders AS
SELECT f.*, c.segment
FROM edw.fact_orders f
JOIN edw.dim_customer c ON c.cust_id = f.cust_id;
-- SCD2 row:
-- cust_sk, cust_id, name, start_dt, end_dt, is_current
```

5. Recursive views without a depth cap can spool out. Hard RI on a 2-billion row fact can destroy load windows.

199. What is a locking view?

1. Definition you should say first: A view that includes LOCKING modifiers. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

200. LOCKING ROW FOR ACCESS?

1. Definition you should say first: Allows dirty reads to reduce blocking. Then spend the rest of the answer on architecture,

JobKit - Teradata DBA interview questions

a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

201. When is ACCESS locking dangerous?

1. Definition you should say first: When the reader must see committed consistent numbers. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

202. What is a deadlock victim?

1. Definition you should say first: The session Teradata aborts to break a deadlock. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Locks: row hash, table, database. ACCESS locking reads dirty. Deadlock: two sessions wait; one is victim. Dummy row lock shows up on some dictionary ops.

3. Use LOCKING ROW FOR ACCESS on reports that can stand dirty reads. For finance numbers, READ and a proper window. Viewpoint blocked sessions shows the blocker.

4. Working code / syntax (write this on Teradata SQL):

```
LOCKING ROW FOR ACCESS
SELECT SUM(amt) FROM edw.fact_orders WHERE order_dt = CURRENT_DATE;
-- find blockers via Viewpoint or DBC.LockInfoV / SessionInfo
ABORT SESSION hostid, sessionid;
```

5. ACCESS during a load can show half-committed totals to a CFO dashboard. Know when that is unacceptable.

203. How do you tune a product join?

1. Definition you should say first: Stats, better join keys, or rewrite to merge/hash. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. EXPLAIN is the plan. Confidence is estimate quality. Merge/hash are good when keys match. Product join is cartesian-ish. Redistribute vs duplicate is how join geography is built. Residual is leftover filter.

3. Fix product joins with stats, better ON keys, or rewrite. Look for 'duplicating all rows of a 2-billion row table'.

4. Working code / syntax (write this on Teradata SQL):

```
EXPLAIN
SELECT f.order_id, c.segment
FROM edw.fact_orders f
INNER JOIN edw.dim_customer c ON c.cust_id = f.cust_id
WHERE f.order_dt BETWEEN DATE '2026-08-01' AND DATE '2026-08-31';
```

5. SET tables with duplicate checks plus a product join will run until spool dies.

204. What is a character set issue?

1. Definition you should say first: Latin vs Unicode; conversions can cost CPU. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table

JobKit - Teradata DBA interview questions

or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
.QUIT 0;
.LABEL FAIL;
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

205. What is a case-specific compare?

1. Definition you should say first: CS vs CASESPECIFIC in comparisons. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
.QUIT 0;
.LABEL FAIL;
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

206. What is a format in BTEQ?

1. Definition you should say first: How numbers and dates print in reports. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
.QUIT 0;
.LABEL FAIL;
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

207. What is .IF ERRORCODE in BTEQ?

1. Definition you should say first: Branch on failure in a script. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;
.IF ERRORCODE <> 0 THEN .GOTO FAIL;
.QUIT 0;
```

JobKit - Teradata DBA interview questions

```
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

208. What is a return code in jobs?

1. Definition you should say first: Scheduler uses it to fail the box. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;  
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
.QUIT 0;  
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

209. How do you alert on a failed load?

1. Definition you should say first: Scheduler plus Viewpoint alerts plus mail. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;  
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
.QUIT 0;  
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

210. What is a runbook?

1. Definition you should say first: Step-by-step for a repeat incident. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;  
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
.QUIT 0;  
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

211. What is a CAB?

1. Definition you should say first: Change advisory board for production changes. Then spend the rest of the answer on architecture, a working example, and a production failure.

JobKit - Teradata DBA interview questions

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;  
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
.QUIT 0;  
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.

212. What is a back-out plan?

1. Definition you should say first: How you undo a change if it fails. Then spend the rest of the answer on architecture, a working example, and a production failure.

2. Latin vs Unicode conversions cost CPU. CS comparisons surprise people. BTEQ .IF ERRORCODE drives job return codes the scheduler uses. CAB and back-out are change control.

3. Every load job: IF ERRORCODE <> 0 then .GOTO FAIL and non-zero exit. Alert Viewpoint + mail. Back-out = restore table or rerun from stg.

4. Working code / syntax (write this on Teradata SQL):

```
.LOGON td/user,pass;  
DELETE FROM edw.fact_orders WHERE order_dt = DATE '2026-09-03';  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
INSERT INTO edw.fact_orders SELECT * FROM stg.orders_in;  
.IF ERRORCODE <> 0 THEN .GOTO FAIL;  
.QUIT 0;  
.LABEL FAIL;  
.QUIT 12;
```

5. A job that always exits 0 even on FastLoad fail will ship empty cubes to finance.